

Homework 2: Empirical Risk and Simple Linear Regression

EECS 245, Spring 2026 at the University of Michigan

due Wednesday, May 13th, 2026 at 11:59PM Ann Arbor Time

Write your solutions to the following problems either by writing them on a piece of paper or on a tablet and scanning your answers as a PDF. Note that you are not allowed to use LaTeX, Google Docs, or any other digital document creation software to type your answers. Homeworks are due to Gradescope by 11:59PM on the due date. See the [syllabus](#) for details on the slip day policy.

Homework will be evaluated not only on the correctness of your answers, but on your ability to present your ideas clearly and logically. You should always explain and justify your conclusions, using sound reasoning. Your goal should be to convince the reader of your assertions. If a question does not require explanation, it will be explicitly stated.

Before proceeding, make sure you're familiar with the [collaboration policy](#).

Total Points: $10 + 7 + 9 + 8 + 9 + 9 = 52$

Problem 1: Homework 1 Solutions Review (10 pts)

Review the solutions to Homework 1. Pick **two problem parts** (for example, Problem 3a and Problem 5) from Homework 1 in which your solutions have the most room for improvement, i.e., where they have unsound reasoning, could be significantly more efficient or clearer, etc. **Include a screenshot of your solution to each problem part**, and in a few sentences, explain what was deficient and how it could be fixed.

Alternatively, if you think one of your solutions is significantly better than the posted one, copy it here and explain why you think it is better. If you didn't do Homework 1, choose two problem parts from it that look challenging to you, and in a few sentences, explain the key ideas behind their solutions in your own words.

Problem 2: Stonks (7 pts)

This problem will eventually have something to do with machine learning. But first, a life lesson.

Suppose you invest in a stock, and:

- In year 1, your investment increases by 50%.
- In year 2, your investment decreases by 50%.
- In year 3, your investment increases by 50%.
- In year 4, your investment decreases by 50%.

What is the average growth rate of your investment, **per year**? The answer is not 0%, because ultimately you've **lost money**, even though it looks like the gains and losses should cancel out.

Why? At end of year 1, you have more money than you started with, and so losing 50% of that money in year 2 hurts more than losing 50% of your starting amount. Then, going up 50% in year 3 earns you less money than originally going up 50% in year 1 did, and so on.

Before we calculate the average growth rate, let's calculate the final value of your investment. To do so, we should convert these growth rates from percentages multipliers, using the formula:

$$\text{multiplier} = 1 + \frac{\text{percentage}}{100}$$

So,

$$\text{final value} = \text{initial value} \cdot \underbrace{1.5}_{\text{year 1}} \cdot \underbrace{0.5}_{\text{year 2}} \cdot \underbrace{1.5}_{\text{year 3}} \cdot \underbrace{0.5}_{\text{year 4}} = \text{initial value} \cdot 0.5625$$

Converting 0.5625 from a percentage back to a growth rate gives us:

$$\text{percentage} = 0.5625 - 1 = -0.4375 = -43.75\%$$

So, in total, we lost 43.75% of our money.

This doesn't give us our average growth rate, though. The average growth rate, as a multiplier, should be a value g such that if our investment grows by g each year, we end up with $1.5 \cdot 0.5 \cdot 1.5 \cdot 0.5 \cdot \text{initial value}$. In other words:

$$\text{final value} = \text{initial value} \cdot g^4$$

So, as a multiplier, we have that:

$$g^4 = 1.5 \cdot 0.5 \cdot 1.5 \cdot 0.5 \implies g = (1.5 \cdot 0.5 \cdot 1.5 \cdot 0.5)^{1/4} \approx 0.8660$$

Converting g back to a percentage gives us:

$$\text{percentage} = 0.8660 - 1 = -0.1340 = -13.40\%$$

So, the average growth rate of our investment, **per year**, is -13.40% — not the 0% that we might initially guess.

What does this have to do with machine learning? Let's re-visit one particular calculation above.

$$g = (1.5 \cdot 0.5 \cdot 1.5 \cdot 0.5)^{1/4}$$

Here, g , is the **geometric mean** of the numbers 1.5, 0.5, 1.5, and 0.5. Geometric means are useful in computing the average of growth rates (when expressed as multipliers). In general, if $y_1, y_2, \dots, y_n$ are **positive** numbers, then their geometric mean is:

$$(y_1 \cdot y_2 \cdot \dots \cdot y_n)^{1/n} = \left(\prod_{i=1}^n y_i \right)^{1/n}$$

Like the arithmetic mean, as we saw in [Chapter 1.2](#), and the harmonic mean, as we saw in Lab 2, the geometric mean is the constant prediction that minimizes average loss for some loss function.

In this case, the loss function is the log-quotient loss, defined as:

$$L_{LQ}(y_i, h(x_i)) = \left[\log \left(\frac{y_i}{h(x_i)} \right) \right]^2$$

Note that $\log(\cdot)$ is the natural logarithm, with base e .

Prove that the geometric mean of $y_1, y_2, \dots, y_n$ is the constant prediction that minimizes average log-quotient loss for the constant model, i.e. that the geometric mean minimizes:

$$R_{LQ}(w) = \frac{1}{n} \sum_{i=1}^n \left[\log \left(\frac{y_i}{w} \right) \right]^2$$

Hint: This is a question involving the three-step modeling process. You'll want to start by finding $\frac{d}{dw} R_{LQ}(w)$ and setting that to 0. As a sub-problem, you'll need to find $\frac{d}{dw} \left[\log \left(\frac{y_i}{w} \right) \right]$. Work one step at a time and make sure your logic is clearly justified. Review the logarithm rules presented in [Homework 1, Problem 5](#), and also use the fact that if $b = \log(a)$, then $a = e^b$.

Problem 3: Slippery Slope (9 pts)

In [Chapter 1.3](#), we found that $w^* = \text{Median}(y_1, y_2, \dots, y_n)$ is the constant prediction that minimizes mean absolute error:

$$R_{\text{abs}}(w) = \frac{1}{n} \sum_{i=1}^n |y_i - w|$$

Suppose that we have a dataset of numbers $y_1, y_2, \dots, y_n$ such that n is **odd** and the values are arranged in increasing order. That is, $y_1 \leq y_2 \leq \dots \leq y_n$.

Note: Parts a) and b) are independent of each other.

- a) (5 pts) Suppose that $R_{\text{abs}}(\alpha) = V$, where V is the minimum value of $R_{\text{abs}}(w)$ and α is one of the numbers in our dataset.

Let $\alpha + \beta$ be the smallest value greater than α in our dataset, where $\beta > 0$. Another way of thinking about this is that $\beta = (\text{smallest value greater than } \alpha) - \alpha$.

Suppose we modify our dataset by replacing the value α with the value $\alpha + \beta + 1$. In our **new** dataset of n values:

- (i) What value of w minimizes $R_{\text{abs}}(w)$?
- (ii) What is the new minimum value of $R_{\text{abs}}(w)$?

Both of your answers should be expressions involving V , α , β , and/or constants.

Hint: Think about the problem on your own for a while. If you're stuck, watch this [hint video](#).

- b) (4 pts) Let y_a and y_b be two values in our dataset such that $y_a < y_b$ and that the slope of $R_{\text{abs}}(w)$ between $w = y_a$ and $w = y_b$ is constant, and equal to $-\frac{2}{3}$.

Suppose we introduce a new value to our dataset that is less than y_a . In our **new** dataset of $n + 1$ values, what is the slope of $R_{\text{abs}}(w)$ between $w = y_a$ and $w = y_b$? Your answer should be an expression involving n and/or constants, but should not contain a or b , or any value of y .

Problem 4: Fun with Correlation (8 pts)

As we will see in [Chapter 2.4](#), the correlation coefficient r between two variables x and y measures the strength of the linear association between them, or intuitively, how tightly the points cluster around a line. Formally, r is defined as:

$$r = \frac{1}{n} \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{\sigma_x} \right) \left(\frac{y_i - \bar{y}}{\sigma_y} \right)$$

where $\bar{x}$ and $\bar{y}$ are the means of x and y , respectively, and σ_x and σ_y are the standard deviations of x and y , respectively.

- a) (3 pts) Let r be the correlation coefficient between x and y . Let t be a new variable defined as:

$$t_i = -2x_i + 5, \quad i = 1, \dots, n$$

Let r' be the correlation coefficient between t and y . Prove that $r' = -r$.

Hint: You can use the facts that if $t_i = ax_i + b$, then $\bar{t} = a\bar{x} + b$ and $\sigma_t = |a|\sigma_x$, without proof. Everything else must be derived from the definition of the correlation coefficient.

- b) (5 pts) Suppose we fit two simple linear regression models by minimizing mean squared error.

- Model 1: predicted $y_i = h(x_i) = w_0^* + w_1^*x_i$
- Model 2: predicted $y_i = h'(t_i) = w_0' + w_1't_i$

(The ' does not indicate a derivative here!)

We already know that $r' = -r$. How do the other quantities compare between the two lines?

- Express w_1' in terms of w_1^* , w_0^* , and/or constants (but no other variables).
- Express w_0' in terms of w_0^* , w_1^* , and/or constants (but no other variables).
- Above, you should have found that the new slope, w_1' , and new intercept, w_0' , are different than the original slope and intercept. However, it turns out that the mean squared error of both model's predictions are the same. That is:

$$\frac{1}{n} \sum_{i=1}^n (y_i - (w_0' + w_1't_i))^2 = \frac{1}{n} \sum_{i=1}^n (y_i - (w_0^* + w_1^*x_i))^2$$

Give a two-sentence English explanation of why this is the case.

- c) (0 pts, **optional**) This part is challenging and potentially time-consuming, so we've made it optional. It's good exam practice though, so if you don't do it now, you should return to it later on when you have more time. It is independent of the previous two parts of this problem.

Prove that, for any dataset $(x_1, y_1), \dots, (x_n, y_n)$ with a correlation coefficient r ,

$$\underbrace{\frac{1}{n} \sum_{i=1}^n (y_i - (w_0^* + w_1^*x_i))^2}_{\text{mean squared error of optimal SLR model}} = \underbrace{\sigma_y^2(1 - r^2)}_{\text{function of correlation coefficient}}$$

Problem 5: Switching Sides (9 pts)

Consider two datasets, A and B . Both datasets have $n = 50$ points, of which 49 are identical, and only one is different between the two datasets:

- **Dataset A :** $(22, 10), (x_2, y_2), \dots, (x_{49}, y_{49}), (x_{50}, y_{50})$
- **Dataset B :** $(22, 50), \underbrace{(x_2, y_2), \dots, (x_{49}, y_{49}), (x_{50}, y_{50})}_{\text{identical in both datasets}}$

Suppose that in both datasets, the x -values have a mean of $\bar{x} = 26$ and a standard deviation of $\sigma_x = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2} = 3$.

- a) (5 pts) Suppose we fit a simple linear regression model by minimizing mean squared error, separately for each dataset.

Let w_1^A and w_1^B be the optimal slopes for datasets A and B , respectively. Determine the difference between w_1^B and w_1^A . That is, find:

$$w_1^B - w_1^A$$

Your answer should be a number with no variables.

Hint: There are many equivalent formulas for the slope of the regression line. We recommend using this one for this problem:

$$w_1^* = \frac{\sum_{i=1}^n (x_i - \bar{x}) y_i}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

- b) (4 pts) Let h_A and h_B be the simple linear regression lines for datasets A and B , respectively. That is, $h_A(x_i) = w_0^A + w_1^A x_i$ and $h_B(x_i) = w_0^B + w_1^B x_i$.

Which of the following values is greater: $|h_A(40) - h_B(40)|$ or $|h_A(42) - h_B(42)|$? Why?

Hint: Intuitively, we're asking which input's predicted value changes more by switching from A to B . Don't try and expand the absolute differences or find their values exactly. Instead, draw a picture of both lines. For each line, there is one point that it is guaranteed to pass through. Using your knowledge of that point, and the slopes of the lines, you should be able to reason about which difference is greater. In your picture, assume that the lines intersect at some point to the left of $x = 40$.

Problem 6: Simple LAD (9 pts)

This problem involves writing code and submitting it to the Gradescope autograder.

There are two ways to access the supplemental Jupyter Notebook:

- **Option 1 (preferred):** Set up a Jupyter Notebook environment locally, use `git` to clone our course repository, and open `homeworks/hw02/hw02.ipynb`. For instructions on how to do this, see the [Environment Setup](#) page of the course website.
- **Option 2:** Click [here](#) to open `hw02.ipynb` on DataHub. Before doing so, read the instructions on the [Environment Setup](#) page on how to use the DataHub.

To receive credit for the programming portion of the homework, you'll need to submit your completed notebook to the autograder on Gradescope. Your submission time for Homework 2 is the **latter** of your PDF and code submission times.